

Approximate Clustering Algorithms for High Dimensional Streaming and Distributed Data

A Dissertation Defense by:
Lee A. Carraher

Department of Electrical Engineering and Computing Systems
University of Cincinnati

November 12, 2017

- ▶ Classification Problem
- ▶ Data Analysis and Clustering
- ▶ Research Goals Motivation and Hypothesis
- ▶ Related Work and Background
- ▶ RPHash and Streaming RPHash Algorithms
- ▶ Experimental Setup, First comparisons
- ▶ Adaptive LSH and Tree-Walk Improvement
- ▶ Experiments with TWRP
- ▶ Conclusions and Future Directions


- ▶ Animals evolved to classify
- ▶ classification tasks for survival
 - ▶ edible, predator, suitable mate
- ▶ Can't experience everything
- ▶ Instead need models based on attributes
 - ▶ teeth size?, leaf shape?, is an engineer?
- ▶ Humans are good classifiers for some things



We are good at identifying people.



Not good at large global networking tasks.



Or abstract data clustering tasks.

- ▶ We are good at classify things we experience
- ▶ Not good at classifying numerical data
- ▶ Not good at classifying large data
 - ▶ # groups, # attributes
- ▶ Not good at classifying unfamiliar things (EKG, Subatomic Interactions...)


A researcher wants to track the geographic dependence or independence, of a particular medical condition. However the patient data reside in disconnected, separately managed databases. The data contains a wide variety of symptoms and biometric measures, for millions of patients. There is an additional constrain placed on the data due to its private and personally identifying nature.

- ▶ grouping the symptoms and referencing location could prove or disprove this claim
- ▶ high dimensional data
- ▶ geographically separated data
- ▶ private individually identifiable data

- ▶ Data Analysis can provide insights into data that often go beyond human cognition
 - ▶ Data with complex relationships (financial trans.)
 - ▶ Things foreign to us (subatomic, galactic objects)
 - ▶ Lots of fast moving things (tweets)
- ▶ How do we create models for these things?

- ▶ Data clustering builds models for classification
- ▶ Cluster Models capture similarities between observation attributes
- ▶ Clustering is a classic Machine Learning Problem
- ▶ The k -Means Problem is a clustering problem

4-means



- ▶ partition observations into k subsets
- ▶ group observations with similar attributes
- ▶ an optimization problem: maximize some objective either inter-cluster or intra-cluster


Theorem (Clustering Objective Function)

$$\operatorname{argmin}_C \sum_{i=1}^k \sum_{x \in C_i} \|x - \mu_i\|^2$$

- ▶ ***k*-Means** \in **Max-SNP** \in **APX** \in **NP**,
- ▶ **APX** are approximable by **PTAS**

Data sets are growing fast

- ▶ Medical, Scientific, Financial, Social, Security



- ▶ Some are ∞ (unbounded) \rightarrow streaming data model.

What needs to be solved?

- ▶ PTAS is good complexity for small problems
- ▶ Big Data problems require lower complexity
- ▶ Streaming algorithms put additional requirements on space
- ▶ If we distribute data, are the connections trustworthy?
 - ▶ Need security

- ▶ A distributable k -means clustering algorithm.
- ▶ A streaming k -means clustering algorithm.
- ▶ Compare clustering performance to other algorithms.
- ▶ Show a sub-quadratic complexity growth.
- ▶ Evaluate dataset security.
- ▶ Experiment with approximate components of RPHash

* same as the proposal

RPHash is a degenerate case of Locality Sensitive Hash (LSH) based *cr*-NN

- ▶ *cr*-NN: given a query x return the k nearest neighbors

Definition (Nearest Neighbor)

[?]

Given a set of vectors P in R^d and query vector q return k vectors $p \subseteq P$ such that $p = \text{Argmin}\{\text{dist}(p', q)\}$, where dist is some metric function.

Build the DB:

```
forall the  $x \in X$  do  
  |  $\text{id} = \text{LSH\_Hash}(x)$   
  |  $T[\text{id}].\text{add}(x)$   
end  
Return:  $T$ 
```

Query the DB:

```
 $\text{id} = \text{LSH\_Hash}(q)$   
Return:  $\text{linear\_scan}(T[\text{id}], k)$ 
```

- ▶ LSH NN: LSH to narrow search
- ▶ Works well in practice!
- ▶ Not so well if LSH buckets contain many points
- ▶ But can be viewed as centroid candidates

- ▶ Degenerate LSH is useful for Clustering
- ▶ Combine Random Projection with LSH gives us an approximate dense region sampling
- ▶ Approximate Clustering is equivalent to local minima clustering

- ▶ Standard algorithm classes
 - ▶ k -means(Lloyd), mean-shift, agglomerative, spectral
- ▶ tend to have bad scalability complexity
- ▶ Focus on k -means(Lloyd) type
- ▶ Tree Based Clustering [?]

- ▶ Projection based clustering
 - ▶ Proclus - project to lower dim.
 - ▶ Cluster Ensemble - Histograms
- ▶ Density scanning clustering algorithms
 - ▶ DBScan
 - ▶ Clique
 - ▶ CLARANS

- ▶ CSketch - most similar
- ▶ Streaming k -Means - similar structure
- ▶ Damped Sliding Window
- ▶ DStream
- ▶ Biased Reservoir Sampling

In this thesis we present 3 algorithms for large scale distributed data clustering, that stem from the same basic setup.

- ▶ Random Projection Hash (RPHash)
- ▶ Streaming RPHash
- ▶ Tree-Walk RPHash (TWRP)

- ▶ Random Projection
- ▶ Locality Sensitive Hashing
- ▶ Exact and approximate counting
- ▶ Off-line clustering

- ▶ Use Random Projection to mitigate Curse of Dimensionality and enforce data embeddings
- ▶ JL-Lemma $d \approx \Omega \left(\frac{\log(m)}{\epsilon^2 \log(1/\epsilon)} \right)$
- ▶ d can be even lower for clustering tasks (Bartal '11)
- ▶ bonus, unrecoverable data anonymization (Liu, Kargupta, Ryan '06)
- ▶ Gaussian RP, DB friendly, and FJLT Projection

Definition (Locality Sensitive Hash Function)

let $\mathbb{H} = \{h : S \rightarrow U\}$ is (r_1, r_2, p_1, p_2) -sensitive if for any $u, v \in S$

1. if $d(u, v) \leq r_1$ then $Pr_{\mathbb{H}}[h(u) = h(v)] \geq p_1$
2. if $d(u, v) > r_2$ then $Pr_{\mathbb{H}}[h(u) = h(v)] \leq p_2$

- ▶ We increase selectivity by appending LSH function results
- ▶ And decrease prob. of missing NN by probing many times


Lattice Based LSH

- ▶ Leech Lattice (Λ_{24}) - optimal regular lattice partitioning in \mathbb{R}^{24}
- ▶ E_8 - optimal regular lattice partitioning in \mathbb{R}^8
- ▶ D_n - checkerboard lattice

Other LSH

- ▶ Spherical - force data to lie on the partitioned surface of a hypersphere
- ▶ P-stable - partition orthogonal projections of a vector with a stable distribution
- ▶ k -means - use the dual of k -means centroids to define the space partitioning

LSH Based Clustering

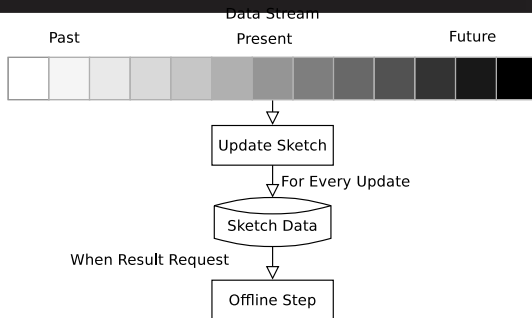


- ▶ k - number of clusters
- ▶ $X = \{x_1, \dots, x_n\}$, $x_k \in \mathbb{R}^m$ - set of data vectors
- ▶ $\mathbb{H}(\cdot)$ - LSH Function with bucket radius= r , dim= d
- ▶ $\mathbb{P} = \{p_1, \dots, p_n\}$ - set of $n, m \times d$ matrices w/ JL property
- ▶ $C : h \rightarrow 0$ - empty map of hashes to counts
- ▶ $M : h \rightarrow \langle 0, \dots, 0 \rangle$ - empty map of hashes to vectors
- ▶ $\text{clusterer}(M, k)$ - a standard clustering algorithm

```
forall the  $x_k \in X$  do
  forall the  $p_i \in \mathbb{P}$  do
     $\tilde{x}_k \leftarrow \sqrt{\frac{m}{d}} p_i^\top x_k$ 
     $t = \mathbb{H}(\tilde{x}_k)$ 
     $C[h]_+ = 1$ 
  end
end
C.sort()
Result:  $C[0 : k * P.length]$ 
```


```
forall the  $x_k \in X$  do
  forall the  $p_i \in \mathbb{P}$  do
     $\tilde{x}_k \leftarrow \sqrt{\frac{m}{d}} p_i^\top x_k$ 
     $h = \mathbb{H}(\tilde{x}_k)$ 
    if  $h \cap C.keys = \emptyset$  then
       $\Delta = M[h] - x_k M[h] = M[h] + \Delta / C[h]$ 
    end
  end
end
end
Result: clusterer( $M, k$ ) ;           // merge centroids
```

Streaming Model




- ▶ Data arrives as a stream
- ▶ No random access to all data vectors
- ▶ data stream is unbounded
 - ▶ algorithm must be sub-quadratic
- ▶ Strict Memory bound

CountMin Sketch Structure



Streaming RPHash Diagram



Data

- ▶ k - number of clusters
- ▶ $x \in \mathbb{R}^m$ - data vector from stream
- ▶ $\mathbb{H}(\cdot)$ - LSH Function radius= r , dim= d
- ▶ $\mathbb{P} = m \times d$ matrix w/ JL property
- ▶ M - lsh_key \rightarrow centroid map
- ▶ C - cm-sketch data structure
- ▶ T - CM-Sketch based ϵk bounded priority queue

forall the $x \in X$ **do**

$$\tilde{x} := \sqrt{\frac{m}{d}} p^T x$$

$t := \mathbb{H}(\tilde{x})$ $C.add(t)$

if $t \in M.keys$ **then**

$M[t].wadd(x)$

else

$M[t] = \text{new centroid}(x, C.count(t))$

$T.insert(M[t])$

end

$M.remove(T.pop())$

end

- ▶ Define evaluation metrics
- ▶ Review Datasets
 - ▶ Real World
 - ▶ Synthetic
- ▶ Optimize RPHash Configuration Parameters
- ▶ Algorithms for Comparison
- ▶ Standard RPHash Results
- ▶ Streaming Results

- ▶ Internal performance metrics - unlabeled metrics
 - ▶ WCSSE - within cluster sum of squares
- ▶ External performance metrics - require ground truth labels
 - ▶ ARI - Adjusted Rand Index
 - ▶ Cluster Purity
- ▶ Runtime
- ▶ Memory Consumption


- ▶ Synthetic data generators in R and Java
 - ▶ Uniformly Distributed Models
 - ▶ Uneven sized clusters
 - ▶ Gaussian distributed vectors
 - ▶ Sparse
 - ▶ Noise injected

| Data Set | Num of Clusters | Num of Features | Num of Vectors | Type of Data |
|--------------------|-----------------|-----------------|----------------|--------------|
| Arrhythmia [?] | 16 | 279 | 452 | Real |
| CNAE-9 [?] | 9 | 856 | 1080 | Binary |
| Cora [?] | 7 | 1433 | 2708 | Binary |
| Gisette [?] | 2 | 5000 | 7000 | Real |
| Human Activity [?] | 6 | 561 | 10299 | Real |
| UJIIndoorLoc [?] | 3 | 520 | 21000 | Real |
| WebKB [?] | 5 | 1703 | 265 | Binary |

6400 Configurations!

| LSH Algorithm | Projected Dimension(s) | # Projections | # Blurings | Offline Clustering |
|----------------------|------------------------|-----------------|------------|---|
| E_8 | 8 | 1, 2, 3, ..., 8 | 1, 2, 3, 4 | k -means Single Linkage Complete Linkage Average Linkage |
| Multi- E_8 | 8, 16, 24, 32 | | | |
| Leech | 24 | | | |
| Multi-Leech | 24, 48, 72, 96 | | | |
| Lévy p -stable | 8, 16, 24, ..., 80 | | | |
| Cauchy p -stable | | | | |
| Gaussian p -stable | | | | |
| Spherical | | | | |

Best Configurations



| Data Set | Configuration | RPHash | | | | k-means | | | |
|-------------|---------------------------------------|--------|--------|---------|--------|---------|--------|---------|-----------|
| | | ARI | Purity | Runtime | Memory | ARI | Purity | Runtime | Memory |
| Arrhythmia | Multi- $E_9/24/1/2/$ Complete Linkage | 0.2710 | 0.5885 | 0.1277 | 0.1420 | 0.0811 | 0.6069 | 0.5287 | 16.4333 |
| CNAE-9 | Multi-Leech/72/5/2/ k -means | 0.3424 | 0.5707 | 0.7917 | 0.8312 | 0.2798 | 0.5312 | 2.7120 | 165.1167 |
| Cora | Multi-Leech/96/3/3/Complete Linkage | 0.1065 | 0.3988 | 2.6022 | 0.6660 | 0.1158 | 0.4271 | 52.410 | 227.9833 |
| Gisette | Spherical/16/2/3/ k -means | 0.5675 | 0.6218 | 1.7769 | 0.4530 | 0.4610 | 0.6002 | 24.746 | 1485.0667 |
| UJIndoorLoc | Spherical/8/8/2/ k -means | 0.6608 | 0.8268 | 1.3840 | 0.2780 | 0.6954 | 0.7750 | 23.821 | 2850.6500 |
| WebKB | Spherical/16/5/1/ k -means | 0.4210 | 0.7396 | 0.1117 | 0.9015 | 0.4403 | 0.7528 | 0.8087 | 50.1500 |

Static Clustering:

- ▶ **k-Means:** [?].
- ▶ **Agglomerative Hierarchical clustering:** Single, Complete, Average and Ward's method.
- ▶ **Self-organizing Tree Algorithm (SOTA):** [?].
- ▶ **k-Means++** [?]

Streaming Clustering

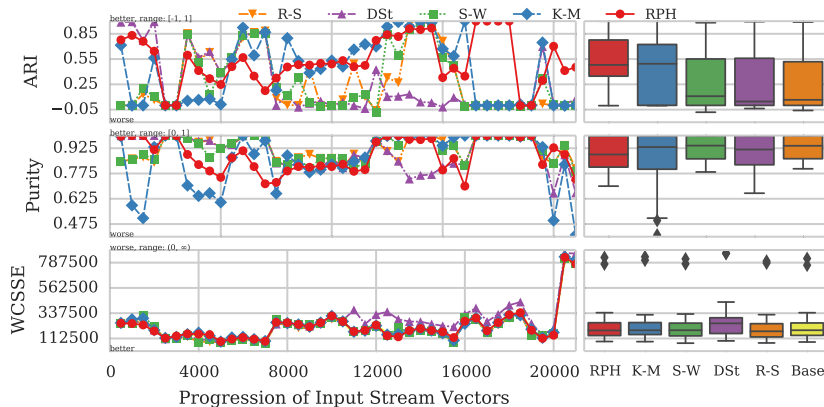
- ▶ **Streaming k-Means:** [?]
- ▶ **Damped Sliding Window:** [?]
- ▶ **DStream:** [?]
- ▶ **Biased Reservoir Sampling:** [?]

Experiments: Real Data

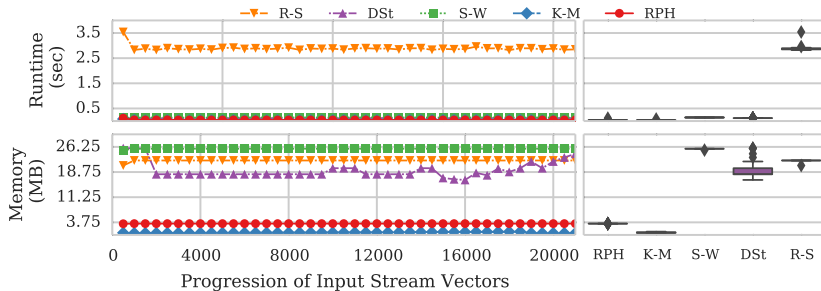
| Data Set | Measures | RPHash | k-means [?] | Single Linkage | Complete Linkage | Average Linkage | Ward's Method [?] | SOTA [?] |
|--------------|----------|--------|-------------|----------------|------------------|-----------------|-------------------|-----------|
| Arrhythmia | ARI | 0.0697 | 0.0811 | 0.0461 | 0.0963 | 0.0546 | 0.0889 | 0.0981 |
| | Purity | 0.6058 | 0.6069 | 0.5730 | 0.5885 | 0.5752 | 0.5951 | 0.6062 |
| | Runtime | 0.2709 | 0.5287 | 0.1680 | 0.1640 | 0.1680 | 0.1680 | 3.4440 |
| | Memory | 0.7070 | 16.4333 | 3.4000 | 3.4000 | 3.4000 | 3.4000 | 21.3000 |
| CNAE-9 | ARI | 0.2788 | 0.2798 | 0.0000 | 0.0000 | 0.0000 | 0.3547 | 0.1730 |
| | Purity | 0.4873 | 0.5312 | 0.1185 | 0.1204 | 0.1185 | 0.5722 | 0.3657 |
| | Runtime | 0.3932 | 2.7120 | 4.3360 | 4.3400 | 4.3400 | 4.3440 | 3.7080 |
| | Memory | 1.1370 | 165.1167 | 24.2000 | 24.2000 | 24.2000 | 24.1000 | 134.2000 |
| Cora | ARI | 0.0915 | 0.1158 | 0.0001 | 0.0120 | 0.0002 | 0.0930 | 0.0647 |
| | Purity | 0.3858 | 0.4271 | 0.3039 | 0.3335 | 0.3039 | 0.4597 | 0.3342 |
| | Runtime | 0.8290 | 52.4100 | 71.9200 | 71.9600 | 71.9400 | 71.9720 | 11.7120 |
| | Memory | 1.4590 | 227.9833 | 100.7000 | 100.7000 | 100.7000 | 100.7000 | 265.4000 |
| Gisette | ARI | 0.1282 | 0.0615 | 0.0000 | 0.0000 | 0.0000 | 0.0018 | 0.1147 |
| | Purity | 0.6720 | 0.6241 | 0.5001 | 0.5003 | 0.5001 | 0.5216 | 0.6694 |
| | Runtime | 2.7363 | 423.7660 | 2280.4320 | 2280.4640 | 2280.0800 | 2280.5480 | 46.8120 |
| | Memory | 1.4300 | 2138.3833 | 829.3000 | 829.3000 | 829.3000 | 829.3000 | 2097.5000 |
| HAR | ARI | 0.3348 | 0.4610 | 0.0000 | 0.3270 | 0.3321 | 0.4909 | 0.3143 |
| | Purity | 0.4631 | 0.6002 | 0.1890 | 0.3770 | 0.3588 | 0.6597 | 0.3966 |
| | Runtime | 1.8774 | 24.7460 | 413.8800 | 414.3320 | 414.0960 | 414.4480 | 14.2440 |
| | Memory | 0.5157 | 1485.0667 | 1259.0000 | 1214.9000 | 1214.8000 | 1214.9000 | 946.2000 |
| UJIIndoorLoc | ARI | 0.5043 | 0.6954 | 0.0001 | 0.0001 | 0.0001 | 0.6021 | 0.3351 |
| | Purity | 0.7105 | 0.7750 | 0.4635 | 0.4635 | 0.4635 | 0.7732 | 0.6918 |
| | Runtime | 2.6363 | 23.8213 | 1093.9440 | 1094.7000 | 1094.8200 | 1095.5360 | 16.1880 |
| | Memory | 0.2460 | 2850.6500 | 5132.4000 | 5049.0000 | 5049.0000 | 5049.0000 | 2227.0000 |
| WebKB | ARI | 0.3205 | 0.4403 | 0.0066 | 0.0404 | 0.0066 | 0.3276 | 0.3906 |
| | Purity | 0.7063 | 0.7528 | 0.4755 | 0.5283 | 0.4755 | 0.7094 | 0.7019 |
| | Runtime | 0.1648 | 0.8087 | 0.3760 | 0.3760 | 0.3760 | 0.3760 | 2.5400 |
| | Memory | 1.2000 | 50.1500 | 6.3000 | 6.4000 | 6.3000 | 6.3000 | 44.1000 |

- ▶ Test on UJI Indoor Localization
- ▶ High Dimensional ($d=561$)
- ▶ Many observation (21000)


Streaming RPHash




Streaming RPHash



Streaming RPHash





- ▶ RPHash and *k*-Means Results are unstable
- ▶ Testing shows that RP and Counting are not the problem
- ▶ Problem must lie in the LSH Function

Fix The LSH Functions


- ▶ LSH functions are good when data is uniformly distributed
- ▶ clusterable data by definition is not uniformly distributed
- ▶ use a set of nested hash functions to adapt to data

Definition (LSH Composability)

An LSH function $\mathbb{H}^n(x)$ that maps $x \in \mathbb{R}^n \rightarrow \mathbb{Z}_2^n$, is composable if there is a related function $\mathbb{H}^{n-1}(x_{n-1})$ that maps $x_{n-1} \in \mathbb{R}^{n-1} \rightarrow \mathbb{Z}_2^{n-1}$ where

$$\mathbb{H}^{n-1}(x_{n-1}) = (\mathbb{H}^n(x) + 1) \cup (\mathbb{H}^n(x) + 0) \text{ for all } x_n \in \mathbb{R}^n$$

Adaptive LSH



Definition (Sign based Projected LSH)

$$H(X) = \sum \text{sign}(P(X))2^n$$

$i = 1$

$ct, ct_prev = C(\mathbb{H}^{i+1}(x)), C(\mathbb{H}^i(x))$


while $i < n$ **and** $2ct > ct_prev$ **do**

$ct_prev, i = ct, i + 1$

$ct = C(\mathbb{H}^i(x))$

end



return $\mathbb{H}^i(x)$




Worse Than Leech and Spherical

- ▶ Composable hashes let us investigate neighbors
- ▶ Use neighbor and parent relationships to decide when cuts are useful
- ▶ Generates a Tree
- ▶ Tree Based Clustering


Tree Walk RPHash




Tree Walk RPHash




Tree Walk RPHash



Tree Walk RPHash




Tree Walk RPHash



- ▶ The tree is exponential in depth
- ▶ $\theta(2^d * m)$ - storage complexity
- ▶ Just need to search, most nodes have low support
- ▶ Count-Min sketch

Count-Min Tree



forall the $x \in X$ **do**

$$\tilde{x} = \sqrt{\frac{m}{d}} p^\top x \quad h := \mathbb{H}(\tilde{x})$$

while $h > 0$ **do**

$$h = h \gg 1$$

$$x' = C[h] + x$$

$$C.\text{add}(h, x')$$

end

end

k -number of clusters

$$X = \{x_1, \dots, x_n\}, x_i \in \mathbb{R}^m$$

C - cm-sketch,

counts \rightarrow vector

\gg - bit shift

$\mathbb{H}(\cdot)$ - LSH Function

$\mathbb{P} = \{p_1, \dots, p_n\}$ - Projectors

$+$ -weighted addition

- ▶ TWRP was developed independently, but similar
- ▶ Concern of CLTree is intersecting clusters
- ▶ We ignore this concern for high dimensional data and proof the following theorem



Theorem (Hyper-rectangle Splitting)


The probability of splitting a hyper-rectangular region into two equal mass clusters where subsequent dimensional cuts are always of the smaller region is 0 as the dimensionality grows to infinity.


$$\lim_{d \rightarrow \infty} \frac{\text{Vol}(R) - \text{Vol}_{\text{removed}}(R)}{\text{Vol}(R)} = 0, \text{ } R \text{ is a hyper-rectangle in } \mathbb{R}^d$$

```
forall the  $H \in \text{sort}(C.ids)$  do  
  | if  $2C[H] < C[H \gg 1]$  then  
    |  $C[H \gg 1] = 0$   
  end  
end  
 $L = []$   
forall the  $h \in \text{sort}(C.counts)$  do  
  |  $L \leftarrow \text{medoid}(C[H])$   
end  
return  $L$ 
```

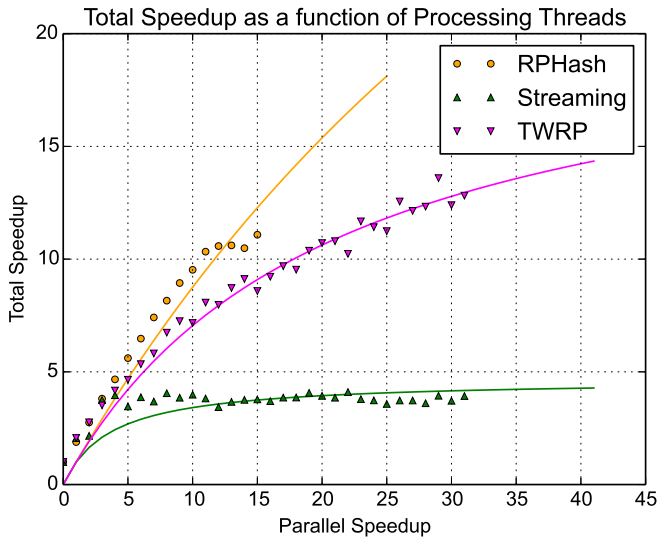
- ▶ Compare Scalability of algorithms
- ▶ Parallel Speedup Comparison
- ▶ Security Evaluation on Real Data







| LSH Algorithm | Time Complexity | Space Complexity |
|----------------------|---------------------------------|---|
| RPHash | $\Theta(nm)$ | $\Theta(nm)$ |
| Streaming RPHash | $\Theta(nm)$ | $\Theta(m \log \log(n))$ |
| TWRP | $\Theta(nm \log^2(n/\epsilon))$ | $\Theta(m \frac{e}{\epsilon} \ln(\frac{n \log(d)}{\sqrt{\delta}}))$ |




- ▶ *RPHash* increases granularity by projection
- ▶ Similar to the *I*-diversity metric
- ▶ Evaluate random projection against data recoverability

$$u = \sqrt{\frac{n}{k}} R_{d \rightarrow s}^T v, v' = \sqrt{\frac{k}{n}} u^T R_{s \rightarrow d}^{-1}$$

- ▶ *R* is non-invertible, best we can do is the Moore-Penrose pseudo-inverse.

Security Assessment (conti.)



- ▶ Empirically we show that TWRP algorithm and both streaming and standard RPHash are comparable to other clustering methods
- ▶ our hypothesis that approximate clustering vs local minima clustering holds for many real world and synthetic datasets
- ▶ RPHash has linear complexity, and memory bound and in the streaming case sub-linear memory bound both in theoretically and shown in experiments.

- ▶ Count-Min Cut Tree is interesting for approximate data analysis
- ▶ Topological Data Analysis could potentially use RPHash for micro-cluster identification to accelerate it
- ▶ Could accelerate hashing with GPUs

Questions??



C. C. Aggarwal.

On biased reservoir sampling in the presence of stream evolution.

In Proceedings of the 32nd International Conference on Very Large Data Bases, VLDB '06, pages 607–618, Seoul, Korea, 2006.



D. Anguita, A. Ghio, L. Oneto, X. Parra, and J. L.

A public domain dataset for human activity recognition using smartphones, Apr. 2013.



D. Arthur and S. Vassilvitskii.

K-means++: The advantages of careful seeding.

In Proceedings of the Eighteenth Annual ACM-SIAM Symposium on Discrete Algorithms, SODA '07, pages

1027–1035, Philadelphia, PA, USA, 2007. Society for Industrial and Applied Mathematics.



V. Braverman, A. Meyerson, R. Ostrovsky, A. Roytman, M. Shindler, and B. Tagiku.

Streaming k-means on well-clusterable data.

In *Proceedings of the Twenty-second Annual ACM-SIAM Symposium on Discrete Algorithms, SODA '11*, pages 26–40. SIAM, 2011.







P. M. Ciarelli and E. Oliveira.

UCI machine learning repository.



H. A. Guvenir.

UCI machine learning repository.

-  I. Guyon, S. R. Gunn, A. Ben-Hur, and G. Dror.
Result analysis of the NIPS 2003 feature selection challenge, 2004.
-  J. A. Hartigan and M. A. Wong.
A k-means clustering algorithm.
JSTOR: Applied Statistics, 28(1):100–108, 1979.
-  J. Herrero, A. Valencia, and J. Dopazo.
A hierarchical unsupervised growing neural network for clustering gene expression patterns, 2001.
-  B. Liu, Y. Xia, and P. S. Yu.
Clustering through decision tree construction.
In Proceedings of the Ninth International Conference on Information and Knowledge Management, CIKM '00, pages 20–29, New York, NY, USA, 2000. ACM.



F. Murtagh and P. Legendre.

Ward's hierarchical agglomerative clustering method:
Which algorithms implement ward's criterion?

J. Classif., 31(3):274–295, Oct. 2014.



H. Samet.

*Foundations of Multidimensional and Metric Data
Structures.*


Morgan Kaufmann, 2006.



P. Sen, G. M. Namata, M. Bilgic, L. Getoor,
B. Gallagher, and T. Eliassi-Rad.

Collective classification in network data.

AI Magazine, 29(3):93–106, 2008.

-  J. Torres-Sospedra, R. Montoliu, A. Martinez-Usó, T. J. Arnau, J. P. Avariento, M. Benedito-Bordonau, and J. Huerta.

Ujiindoorloc: A new multi-building and multi-floor database for wlan fingerprint-based indoor localization problems, 2014.

-  L. Tu and Y. Chen.

Stream data clustering based on grid density and attraction.

ACM Trans. Knowl. Discov. Data, 3(3):12:1–12:27, July 2009.



Y. Zhu and D. Shasha.

Statstream: Statistical monitoring of thousands of data streams in real time.

In Proc of the 28th Int Conf on Very Large Data Bases,
pages 358–369, 2002.